

Manahil Fatima Anwar

20K-0134

BAI-7A

CV Lab 07

Lab Tasks

Task 01

```
In [1]: import cv2
import numpy as np
import matplotlib.pyplot as plt

image = cv2.imread("image01.jpg", cv2.IMREAD_GRAYSCALE)

sobel_x = cv2.Sobel(image, cv2.CV_64F, 1, 0, ksize=5)
sobel_y = cv2.Sobel(image, cv2.CV_64F, 0, 1, ksize=5)
edge_image = cv2.magnitude(sobel_x, sobel_y)

canny_edge_image = cv2.Canny(image, 100, 200)

sigma = 1.5
smoothed_image = cv2.GaussianBlur(image, (0, 0), sigma)
laplacian = cv2.Laplacian(smoothed_image, cv2.CV_64F)

plt.figure(figsize=(12, 6))

plt.subplot(2, 3, 1)
plt.imshow(image, cmap = "gray")
plt.title('Original Image')
plt.axis('off')

plt.subplot(2, 3, 2)
plt.imshow(edge_image, cmap = "gray")
plt.title('Sobel Edge Detection')
plt.axis('off')

plt.subplot(2, 3, 3)
plt.imshow(canny_edge_image, cmap = "gray")
plt.title('Canny Edge Detection')
plt.axis('off')

plt.subplot(2, 3, 4)
plt.imshow(laplacian, cmap = "gray")
plt.title('Laplacian of Gaussian (LoG) Edge Detection')
plt.axis('off')

plt.show()
```

Original Image



Sobel Edge Detection



Canny Edge Detection



Laplacian of Gaussian (LoG) Edge Detection



Task 02

```
In [2]: import numpy as np
import pywt
import pywt.data
import matplotlib.pyplot as plt
import soundfile as sf

audio, sample_rate = sf.read(r"C:\Users\ABC\Desktop\BAI\BAI-S7\CV Lab\Lab 07\St

wavelet = "db4"
level = 5

coeffs = pywt.wavedec(audio, wavelet, level=level)

threshold = 0.1

thresholded_coeffs = [pywt.threshold(c, threshold, mode="soft") for c in coeffs

denoised_audio = pywt.waverec(thresholded_coeffs, wavelet)

sf.write(r"C:\Users\ABC\Desktop\BAI\BAI-S7\CV Lab\Lab 07\clean_audio.mp3", deno

import sounddevice as sd
sd.play(denoised_audio, sample_rate)
sd.wait()

plt.figure(figsize=(10, 6))

plt.subplot(3, 1, 1)
plt.title("Original Audio")
plt.plot(audio)

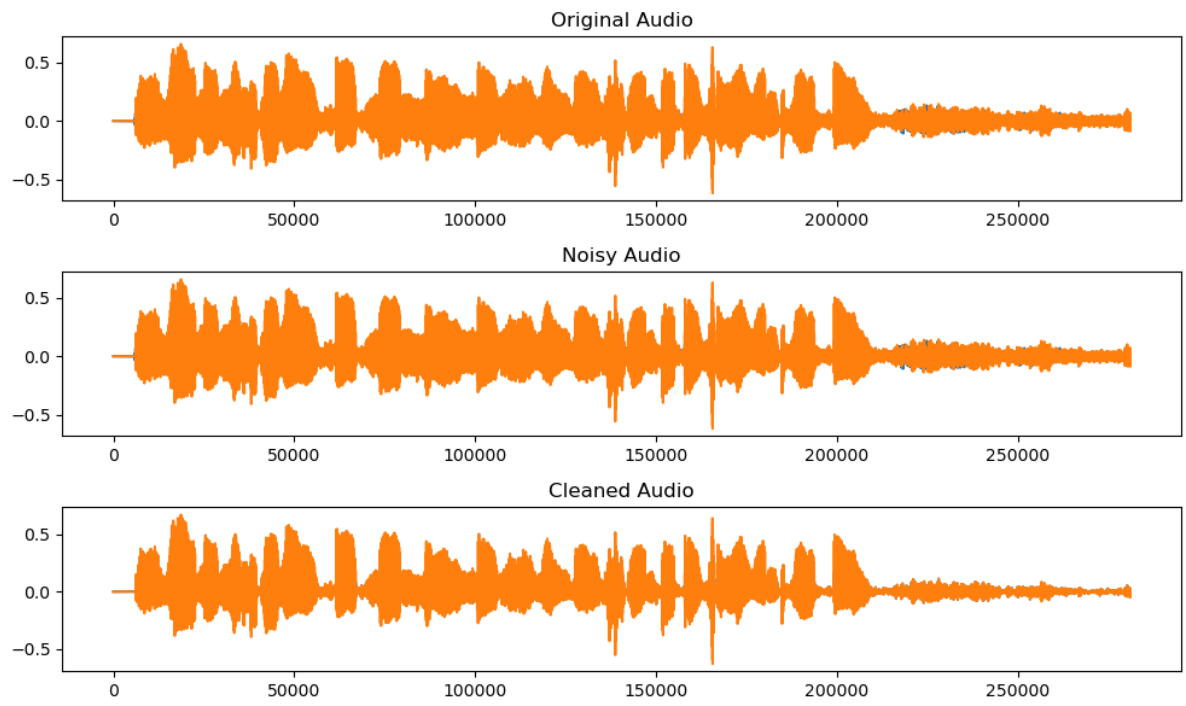
plt.subplot(3, 1, 2)
plt.title("Noisy Audio")
plt.plot(audio)

plt.subplot(3, 1, 3)
plt.title("Cleaned Audio")
plt.plot(denoised_audio)

plt.tight_layout()
plt.show
```

C:\Users\ABC\anaconda3\Lib\site-packages\pywt_multilevel.py:43: UserWarning:
Level value of 5 is too high: all coefficients will experience boundary effects.
warnings.warn(

Out[2]: <function matplotlib.pyplot.show(close=None, block=None)>



Task 03

```
In [3]: img = cv2.imread('image02.jpg', cv2.IMREAD_COLOR)

gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

edges = cv2.Canny(gray, 50, 200)

lines = cv2.HoughLinesP(edges, 0.1, np.pi/180, 80, minLineLength=10, maxLineGap=10)

for line in lines:
    x1, y1, x2, y2 = line[0]
    cv2.line(img, (x1, y1), (x2, y2), (255, 0, 0), 3)

plt.figure(figsize=(12, 6))

plt.subplot(1, 3, 1)
plt.imshow(gray, cmap="gray")
plt.title("Original Image")
plt.axis('off')

plt.subplot(1, 3, 2)
plt.imshow(img, cmap="gray")
plt.title("Lines Detected")
plt.axis('off')

plt.subplot(1, 3, 3)
plt.imshow(edges, cmap="gray")
plt.title("Edges Detected")
plt.axis('off')

plt.show()
```


Original Image



Lines Detected



Edges Detected



Task 04

```
In [4]: image = cv2.imread("image02.jpg", cv2.IMREAD_GRAYSCALE)

sift = cv2.SIFT_create()

keypoints, descriptors = sift.detectAndCompute(image, None)

image_with_keypoints = cv2.drawKeypoints(image, keypoints, None)

plt.imshow(image_with_keypoints)
plt.title("SIFT")
plt.axis('off')

# Match the descriptors with a pre-existing database
# Implement feature matching here using a database of known computer systems ar

# If matches are found, identify and label the computer systems and components
# Implement identification based on the matched features

# Maintain an inventory based on the identified and labeled computer systems ar
# Update the inventory based on the identified computer systems and components
```

Out[4]: (-0.5, 4623.5, 2607.5, -0.5)

SIFT



Task 05


```
In [5]: import numpy as np
import pywt
import matplotlib.pyplot as plt

# Generate a sample sensor dataset
np.random.seed(0)
sensor_data = np.random.normal(0, 1, 1000)

# Choose the wavelet family and decomposition level
# wavelet = "db4"
wavelet = "haar"
level = 3

# Perform the wavelet transformation
coeffs = pywt.wavedec(sensor_data, wavelet, level=level)

# Set a threshold for anomaly detection
threshold = 2.0

# Apply thresholding to the wavelet coefficients
thresholded_coeffs = [pywt.threshold(c, threshold, mode="soft") for c in coeffs]

# Reconstruct the denoised signal
denoised_signal = pywt.waverec(thresholded_coeffs, wavelet)

# Calculate the residuals (differences between original and denoised signal)
residuals = sensor_data - denoised_signal

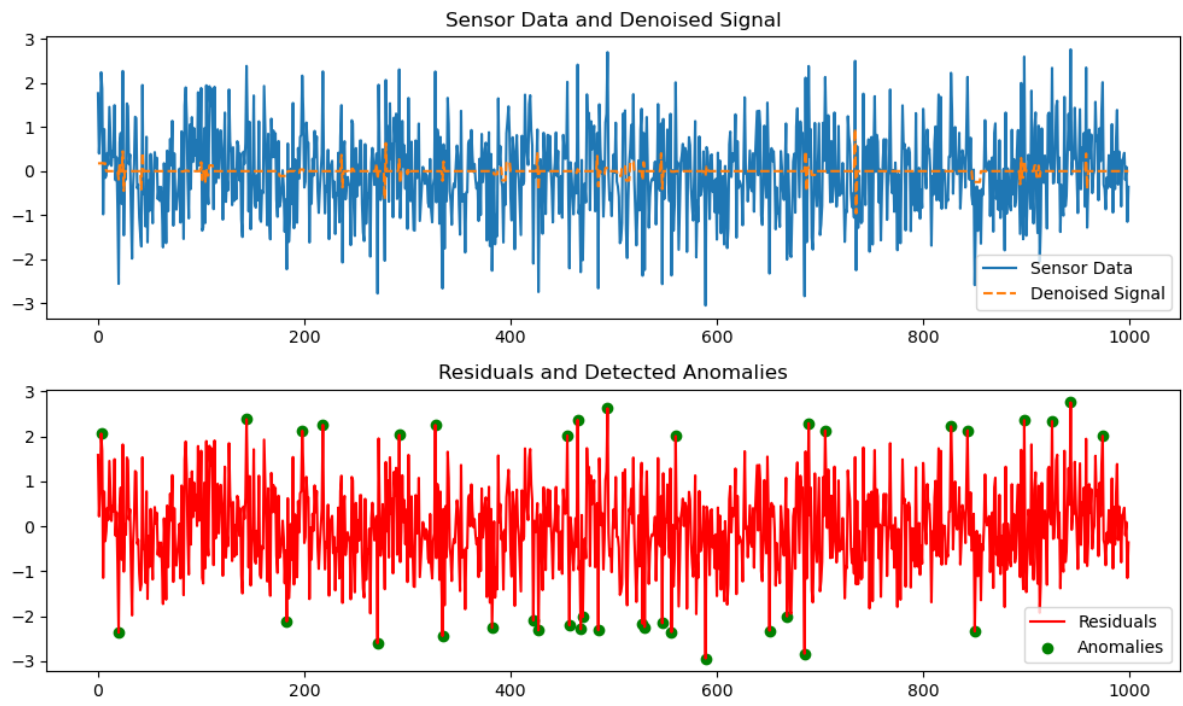
# Detect anomalies based on the residuals and threshold
anomalies = np.where(np.abs(residuals) > threshold)[0]

# Plot the sensor data, denoised signal, and detected anomalies
plt.figure(figsize=(10, 6))
plt.subplot(2, 1, 1)
plt.plot(sensor_data, label="Sensor Data")
plt.plot(denoised_signal, label="Denoised Signal", linestyle="--")
plt.legend()
plt.title("Sensor Data and Denoised Signal")

plt.subplot(2, 1, 2)
plt.plot(residuals, label="Residuals", color="red")
plt.scatter(anomalies, residuals[anomalies], color="green", label="Anomalies")
plt.legend()
plt.title("Residuals and Detected Anomalies")

plt.tight_layout()
plt.show()

# Output detected anomaly indices
print("Detected Anomalies:", anomalies)
```



Home Tasks

Task 01

```

In [6]: import cv2
import numpy as np
import matplotlib.pyplot as plt

def draw_matches(img1, kp1, img2, kp2, matches):
    rows1, cols1 = img1.shape
    rows2, cols2 = img2.shape

    out = np.zeros((max([rows1, rows2]), cols1 + cols2, 3), dtype='uint8')
    out[:rows1, :cols1] = np.dstack([img1, img1, img1])
    out[:rows2, cols1:] = np.dstack([img2, img2, img2])

    for mat in matches:
        img1_idx = mat.queryIdx
        img2_idx = mat.trainIdx
        (x1, y1) = kp1[img1_idx].pt
        (x2, y2) = kp2[img2_idx].pt

        cv2.circle(out, (int(x1), int(y1)), 4, (0, 255, 255), 1)
        cv2.circle(out, (int(x2) + cols1, int(y2)), 4, (0, 255, 255), 1)
        cv2.line(out, (int(x1), int(y1)), (int(x2) + cols1, int(y2)), (0, 255,

    return out

def sift_detector(reference_image, test_images):
    sift = cv2.SIFT_create()

    kp1, des1 = sift.detectAndCompute(reference_image, None)

    fig, axs = plt.subplots(1, len(test_images) + 1, figsize=(15, 6))

    axs[0].imshow(cv2.cvtColor(reference_image, cv2.COLOR_BGR2RGB))
    axs[0].set_title('Reference Image')

    for i, test_image in enumerate(test_images):
        kp2, des2 = sift.detectAndCompute(test_image, None)

        bf = cv2.BFMatcher()
        matches = bf.knnMatch(des1, des2, k=2)

        good_matches = []
        for m, n in matches:
            if m.distance < 0.8 * n.distance:
                good_matches.append(m)

        if len(good_matches) > 5:
            src_pts = np.float32([kp1[m.queryIdx].pt for m in good_matches]).r
            dst_pts = np.float32([kp2[m.trainIdx].pt for m in good_matches]).r

            M, mask = cv2.findHomography(src_pts, dst_pts, cv2.RANSAC, 5.0)
            matchesMask = mask.ravel().tolist()

            h, w = reference_image.shape
            pts = np.float32([[0, 0], [0, h - 1], [w - 1, h - 1], [w - 1, 0]])
            dst = cv2.perspectiveTransform(pts, M)

            test_image = cv2.polylines(test_image, [np.int32(dst)], True, 255,

```

```

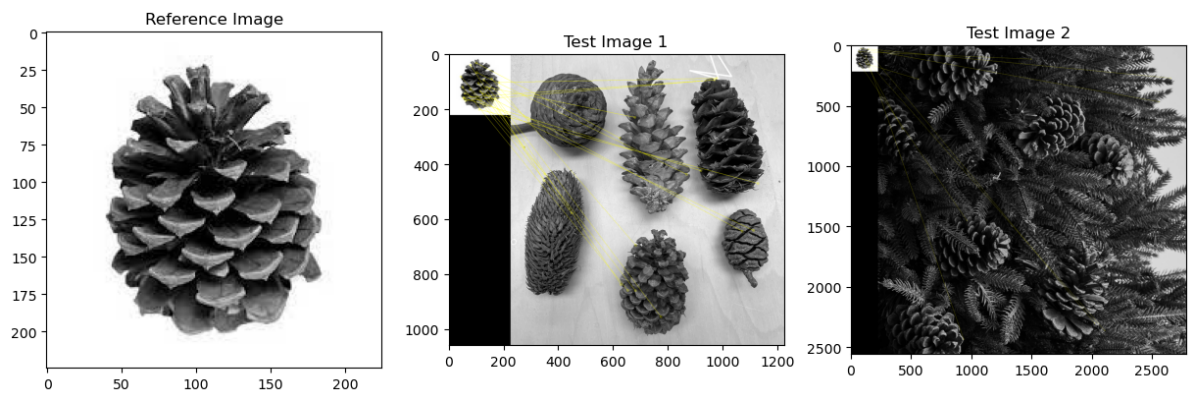
else:
    print("Not enough matches are found for test image {} - {} matches

output_image = draw_matches(reference_image, kp1, test_image, kp2, good
axs[i + 1].imshow(cv2.cvtColor(output_image, cv2.COLOR_BGR2RGB))
axs[i + 1].set_title(f'Test Image {i + 1}')

plt.show()

reference_image = cv2.imread('pine cone_ref.jpg', 0)
test_images = [cv2.imread('pine cone_test01.jpg', 0), cv2.imread('pine cone_test02.jpg', 0)]
sift_detector(reference_image, test_images)

```



Task 02

```
In [18]: import cv2
import numpy as np
from matplotlib import pyplot as plt

# Loading the images
image_filenames = ["my_pic01.jpg", "my_pic02.jpg", "my_pic03.jpg"]
images = [cv2.imread(image) for image in image_filenames]

# Initializing SIFT
sift = cv2.SIFT_create()

# Find keypoints and descriptors for each image
keypoints = []
descriptors = []
for image in images:
    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    kp, des = sift.detectAndCompute(gray, None)
    keypoints.append(kp)
    descriptors.append(des)

# Feature matching between the descriptors of the images
bf = cv2.BFMatcher()
matches = bf.knnMatch(descriptors[0], descriptors[1], k=2)

# Apply ratio test
good = []
for m, n in matches:
    if m.distance < 0.75 * n.distance:
        good.append(m)

# Homography estimation
src_pts = np.float32([keypoints[0][m.queryIdx].pt for m in good]).reshape(-1, 2)
dst_pts = np.float32([keypoints[1][m.trainIdx].pt for m in good]).reshape(-1, 2)

M, mask = cv2.findHomography(src_pts, dst_pts, cv2.RANSAC, 5.0)

# Image stitching
result = cv2.warpPerspective(images[0], M, (images[0].shape[1] + images[1].shape[1], images[0].shape[0]))
result[0:images[1].shape[0], 0:images[1].shape[1]] = images[1]

# Display the panorama
plt.imshow(cv2.cvtColor(result, cv2.COLOR_BGR2RGB))
plt.show()
```



```

In [19]: import cv2
import numpy as np

# List of image file paths containing your images

image_paths = ["my_pic01.jpg", "my_pic02.jpg", "my_pic03.jpg", "my_pic04.jpg"]

# Initialize SIFT
sift = cv2.SIFT_create()

# Initialize the first image
stitched_image = cv2.imread(image_paths[0], cv2.IMREAD_COLOR)

for i in range(1, len(image_paths)):
    # Load the next image
    image2 = cv2.imread(image_paths[i], cv2.IMREAD_COLOR)

    # Detect and compute keypoints and descriptors for both images
    keypoints1, descriptors1 = sift.detectAndCompute(stitched_image, None)
    keypoints2, descriptors2 = sift.detectAndCompute(image2, None)

    # Create a BFMatcher (Brute-Force Matcher)
    bf = cv2.BFMatcher()

    # Match descriptors between the two images
    matches = bf.knnMatch(descriptors1, descriptors2, k=2)

    # Apply ratio test to filter good matches
    good_matches = []
    for m, n in matches:
        if m.distance < 0.75 * n.distance:
            good_matches.append(m)

    # If there are enough good matches, consider it a match
    if len(good_matches) > 10: # You can adjust this threshold
        # Extract matched keypoints' coordinates
        pts1 = np.float32([keypoints1[m.queryIdx].pt for m in good_matches]).reshape((1, -1), order=1)
        pts2 = np.float32([keypoints2[m.trainIdx].pt for m in good_matches]).reshape((1, -1), order=1)

        # Find the perspective transformation (homography) between the two images
        M, mask = cv2.findHomography(pts2, pts1, cv2.RANSAC, 5.0)

        # Warp the second image to align with the first image
        image2_warped = cv2.warpPerspective(image2, M, (stitched_image.shape[1], stitched_image.shape[0]))

        # Append the two images horizontally
        stitched_image = cv2.hconcat([stitched_image, image2_warped])

# Save the final stitched image
cv2.imwrite(r"C:\Users\ABC\Desktop\panoramic_image.jpg", stitched_image)

```



```
-----  
error                                     Traceback (most recent call last)  
Cell In[19], line 19  
    16 image2 = cv2.imread(image_paths[i], cv2.IMREAD_COLOR)  
    18 # Detect and compute keypoints and descriptors for both images  
--> 19 keypoints1, descriptors1 = sift.detectAndCompute(stitched_image, None)  
    20  
    20 keypoints2, descriptors2 = sift.detectAndCompute(image2, None)  
    22 # Create a BFMatcher (Brute-Force Matcher)  
  
error: OpenCV(4.8.1) D:\a\opencv-python\opencv-python\opencv\modules\core\src  
\alloc.cpp:73: error: (-4:Insufficient memory) Failed to allocate 558931968 b  
ytes in function 'cv::OutOfMemoryError'
```

Task 03

```
In [7]: img = cv2.imread('lanes.jpg', cv2.IMREAD_COLOR)

gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

edges = cv2.Canny(gray, 50, 200)

lines = cv2.HoughLinesP(edges, 1, np.pi/180, 68, minLineLength=5, maxLineGap=10)

for line in lines:
    x1, y1, x2, y2 = line[0]
    cv2.line(img, (x1, y1), (x2, y2), (255, 0, 0), 2)

plt.figure(figsize=(12, 6))

plt.subplot(1, 3, 1)
plt.imshow(gray, cmap="gray")
plt.title("Original Image")
plt.axis('off')

plt.subplot(1, 3, 2)
plt.imshow(img, cmap="gray")
plt.title("Lines Detected")
plt.axis('off')

plt.subplot(1, 3, 3)
plt.imshow(edges, cmap="gray")
plt.title("Edges Detected")
plt.axis('off')

plt.show()
```


Original Image



Lines Detected



Edges Detected



Task 04

```
In [8]: img = cv2.imread('coins.jpg', cv2.IMREAD_COLOR)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
img_blur = cv2.medianBlur(gray, 5)
circles = cv2.HoughCircles(img_blur, cv2.HOUGH_GRADIENT, 1, 50, param1 = 100,
if circles is not None:
    circles = np.uint16(np.around(circles))
    for i in circles[0, :]:
        cv2.circle(img, (i[0], i[1]), i[2], (255, 0, 0), 2)
        #cv2.circle(img, (i[0], i[1]), 2, (255, 0, 0), 5)
plt.figure(figsize = (12, 6))
plt.subplot(1, 2, 1)
plt.imshow(gray, cmap='gray')
plt.title('Original')
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(img, cmap='gray')
plt.title('Circles Detected')
plt.axis('off')
plt.show()
```



Task 05

```
In [9]: def detect_boundaries(video_stream, security_zone_coordinates):
        x1, y1, x2, y2 = security_zone_coordinates

        cap = cv2.VideoCapture(video_stream)

        while cap.isOpened():
            ret, frame = cap.read()

            if not ret:
                break

            cv2.rectangle(frame, (x1, y1), (x2, y2), (255, 0, 0), 2)

            gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

            edges = cv2.Canny(gray, 50, 150)

            if np.any(edges[y1:y2, x1:x2] == 255):
                print("Unauthorized object detected!")
                cv2.imshow('Boundary Detection', frame)

                if cv2.waitKey(1) & 0xFF == ord('q'):
                    break
            cap.release()
            cv2.destroyAllWindows()
        security_zone_coordinates = (100, 100, 500, 500)

        video_stream = r'C:\Users\ABC\Videos\Captures\ (3) Example of Hi-Definition Video
        detect_boundaries(video_stream, security_zone_coordinates)
```


Task 06

```
In [10]: def is_screen_on(current_image, reference_image, threshold=0.9):
    gray_current = cv2.cvtColor(current_image, cv2.COLOR_BGR2GRAY)
    gray_reference = cv2.cvtColor(reference_image, cv2.COLOR_BGR2GRAY)

    ssim_index = cv2.matchTemplate(gray_current, gray_reference, cv2.TM_CCOEFF_

    if ssim_index > threshold:
        return True
    else:
        return False

reference_image = cv2.imread('img4.jpg')

current_screen_1 = cv2.imread('img1.jpg')
current_screen_2 = cv2.imread('img2.jpg')
current_screen_3 = cv2.imread('img3.jpg')
current_screen_4 = cv2.imread('img5.jpg')

screens = [current_screen_1, current_screen_2, current_screen_3, current_screen_4]
for i, screen in enumerate(screens, start=1):
    if screen is not None:
        is_on = is_screen_on(screen, reference_image)
        if is_on:
            print(f"The screen {i} is ON")
        else:
            print(f"The screen {i} is OFF")
    else:
        print(f"Failed to capture the current screen {i}.")

plt.figure(figsize = (8, 4))
plt.subplot(1, 5, 1)
plt.imshow(reference_image, cmap='gray')
plt.title('Ref')
plt.axis('off')

plt.subplot(1, 5, 2)
plt.imshow(current_screen_1, cmap='gray')
plt.title('Image 01')
plt.axis('off')

plt.subplot(1, 5, 3)
plt.imshow(current_screen_2, cmap='gray')
plt.title('Image 02')
plt.axis('off')

plt.subplot(1, 5, 4)
plt.imshow(current_screen_3, cmap='gray')
plt.title('Image 03')
plt.axis('off')

plt.subplot(1, 5, 5)
plt.imshow(current_screen_4, cmap='gray')
plt.title('Image 04')
plt.axis('off')
```



```
plt.show()
```

The screen 1 is OFF
 The screen 2 is OFF
 The screen 3 is OFF
 The screen 4 is OFF



Task 07

```
In [17]: reference_image_path = 'paupau img.jpg'
reference_image = cv2.imread(reference_image_path, cv2.IMREAD_GRAYSCALE)

face_cascade = cv2.CascadeClassifier(cv2.data.harcascades + 'haarcascade_frontalface_default.xml')

def detection(video_path):
    cap = cv2.VideoCapture(video_path)

    while True:
        ret, frame = cap.read()

        if not ret:
            break

        gray_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

        faces = face_cascade.detectMultiScale(gray_frame, scaleFactor=1.1, minNeighbors=5, minSize=(30, 30))

        for (x, y, w, h) in faces:
            cv2.rectangle(frame, (x, y), (x + w, y + h), (0, 255, 0), 2)

        cv2.imshow("Frame", frame)
        if cv2.waitKey(1) & 0xFF == ord('q'):
            break

    cap.release()
    cv2.destroyAllWindows()

video_path = 'paupau vid.mp4'
detection(video_path)
```